




STEVENS
INSTITUTE *of* TECHNOLOGY
THE INNOVATION UNIVERSITY®

CPE/EE 695: Applied Machine Learning

Lecture 3-1: Logistic Regression

Dr. Shucheng Yu, Associate Professor
Department of Electrical and Computer Engineering
Stevens Institute of Technology





Logistic Regression

Usually be used for binary **classification**:

$\Pr(Y|X)$, where Y is a binary variable. (why probability?)



Logistic Regression

Usually be used for binary **classification**:

$\Pr(Y|X)$, where Y is a binary variable. (why probability?)

E.g., $\Pr(\text{Tomorrow snow} \mid \text{today windy})$



Logistic Regression

Usually be used for binary **classification**:

$\Pr(Y|X)$, where Y is a binary variable. (why probability?)

E.g., $\Pr(\text{Tomorrow snow} \mid \text{today windy})$

Let $\Pr(Y=1 \mid X=x) = p(x; \theta)$



Logistic Regression

Usually be used for binary **classification**:

$\Pr(Y|X)$, where Y is a binary variable. (why probability?)

E.g., $\Pr(\text{Tomorrow snow} \mid \text{today windy})$

Let $\Pr(Y=1 \mid X=x) = p(x; \theta)$

Maximize likelihood:

$$\prod_{i=1}^n p^{y_i} (1 - p)^{1-y_i}$$



Logistic Regression

Usually be used for binary **classification**:

$\Pr(Y|X)$, where Y is a binary variable. (why probability?)

E.g., $\Pr(\text{Tomorrow snow} \mid \text{today windy})$

Let $\Pr(Y=1 \mid X=x) = p(x; \theta)$

Maximize likelihood:

$$\prod_{i=1}^n p^{y_i} (1 - p)^{1-y_i}$$

Assumption: p is modeled with parameter θ ; otherwise, optimization problem doesn't work



Logistic Regression

Let $\Pr(Y=1 \mid X=x) = p(x; \theta)$

Maximize likelihood:

$$\prod_{i=1}^n p^{y_i} (1 - p)^{1-y_i}$$



Logistic Regression

Let $\Pr(Y=1 \mid X=x) = p(x; \theta)$

Maximize likelihood:

$$\prod_{i=1}^n p^{y_i} (1 - p)^{1-y_i}$$

Task: to estimate θ by maximizing the likelihood.



Logistic Regression

Let $\Pr(Y=1 \mid X=x) = p(x; \theta)$

Maximize likelihood:

$$\prod_{i=1}^n p^{y_i} (1 - p)^{1-y_i}$$

Task: to estimate θ by maximizing the likelihood.

How can we use **linear regression** to solve this?



Logistic Regression

Let $\Pr(Y=1 \mid X=x) = p(x; \theta)$

Task: to estimate θ by maximizing the likelihood.

How can we use **linear regression** to solve this?

Attempt 1: assume $p(x; \theta)$ be a linear function of x

Attempt 2: assume $\log p(x; \theta)$ be a linear function of x

Attempt 3: assume $\log \frac{p}{1-p}$ be a linear function of x (good)

Remember: $0 \leq p \leq 1$



Logistic Regression

Logistic regression model

$$\log \frac{p}{1-p} = \theta_0 + x \cdot \theta$$

Which gives $p(x) = \frac{1}{1 + e^{-\theta_0 + x \cdot \theta}}$

Logistic Regression


Logistic Regression model estimated probability:

$$\hat{p} = h_w(x) = \sigma(\theta^T \cdot x)$$

where $\sigma(\cdot)$ is a logistic function (or sigmoid function).

Prediction:

$$\hat{y} = \begin{cases} 0, & \text{if } \hat{p} < 0.5 \\ 1, & \text{if } \hat{p} \geq 0.5 \end{cases}$$



Training the logistic regression model $\hat{p}(x) = \sigma(\theta^T \cdot x)$ is to learn the best value of parameter θ that makes the model fit the training data.



Logistic Regression

To train a logistic regression model, we first need to define a performance measure. A commonly used measure is so-called the **log loss** function:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)})]$$

It is easier to explain the log loss function with one train example case, in which we want to maximize the posterior probability

$$P(y|x) = \hat{p}^y (1 - \hat{p})^{(1-y)} = \begin{cases} \hat{p}, & \text{when } y = 1 \\ 1 - \hat{p}, & \text{when } y = 0 \end{cases}$$

Take log of both sides, we have $\log P(y|x) = y \log \hat{p} + (1 - y) \log(1 - \hat{p})$.

Average the sum of m training examples, we obtain $J(\theta)$.



Logistic Regression

Learning the logistic regression model is to find:

$$\hat{\theta} = \operatorname{argmin}_{\theta} J(\theta).$$

Where

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)})],$$
$$\hat{p} = \sigma(\theta^T \cdot x)$$

No Normal Equation (i.e., closed form solution) for θ .

But the cost function $J(\theta)$ is **convex** and **derivable**. **Gradient Descent** is guaranteed to find global maximum.



Training Logistic Regression Model

The gradient of the log loss function $J(\theta)$ is:

$$\nabla_{\theta} J(\theta) = \frac{1}{m} \sum_{i=1}^m (\delta(\theta^T \cdot x^{(i)}) - y^{(i)}) x_j^{(i)}$$

At each round of GD, θ is updated as following (similar to linear regression, different values of m for different modes of GD):

$$\theta = \theta - \eta \nabla_{\theta} J(\theta)$$



Training Logistic Regression Model

Overfitting may also happen in logistic regression.

Similarly, to combat overfitting we can introduce a **regularization** term to the cost function $J(\theta)$:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)})] + \lambda R(\theta)$$

where λ is a hyperparameter and $R(\theta)$ can be ℓ_k -norm of θ , i.e.,

$$R(\theta) = \|\theta\|_k = (\sum \theta_i^k)^{\frac{1}{k}}$$

Note: $R(\theta)$ is ℓ_2 -norm in Ridge regression, ℓ_1 -norm in Lasso regression.



Multi-Class Classification

One-Vs-Rest Method:

We can use **binary classifier** for multi-class classification with so-called the One-Vs-Rest (OvR) method. Specifically, it uses multiple rounds of binary classification for multi-class classification.

For example, to determine if an object X is a dog, cat or fish, we call a binary classifier $f()$ as follows:

```
if  $f(X)$  outputs dog
    return dog;
else if  $f(X)$  outputs cat
    return cat;
else return fish
```

Multi-Class Classification

Multinomial Logistic Regression:


Another approach for multi-class classification is to use the multinomial logistic regression. For each class $1 \leq k \leq K$,

1) first compute $z_k(x) = \theta_k^T \cdot x$

2) then compute Softmax function:

$$\hat{p}_k = \delta(z_k(x))_k = \frac{\exp(z_k(x))}{\sum_{j=1}^K \exp(z_j(x))}$$

where θ_k is the vector of parameters of input features for z_k .



Multi-Class Classification

Training Multinomial Logistic Regression Model:

The performance measure is the **cross-entropy** cost function:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} \log(\hat{p}_k^{(i)})$$

where $y_k^{(i)} = \begin{cases} 1, & \text{the } i^{\text{th}} \text{ example of class } k \\ 0, & i^{\text{th}} \text{ example not of class } k \end{cases}$

GD can be used to train the multinomial logistic regression model. The gradient is:

$$\nabla_{\theta_k} J(\theta) = \frac{1}{m} \sum_{i=1}^m (\hat{p}_k^{(i)} - y_k^{(i)}) x^{(i)}$$



Multi-Class Classification

Other Approaches: **One-Vs-One Method:**

The One-Vs-One (OvO) method constructs a **binary classifier** for each pair of classes. Therefore, with **K** classes, we need to construct **$K(K-1)/2$** binary classifiers.

The decision at prediction time can be made by **counting the votes** from individual binary classifiers. In case of a tie, it compares the aggregated classification confidence (i.e., the output probability) of individual binary classifiers of each class and the higher one is selected.

The OvO method is **slower** than OvR. But for some algorithms (e.g., Kernel algorithms) which cannot scale with many training examples, this algorithm can be helpful.

Multi-Class Classification

Other Approaches: Error-Correcting Output Codes

The Error-Correcting Output Codes (ECOC) method encodes **K classes** into **N bit vectors**. Each class is represented as a bit in each bit vector. ECOC trains **N binary classifiers**, each splitting one group of classes from another (using the column bit vectors below). At prediction time, the N binary classifiers are called, the outputs of them yielding an N -bit vector. A class with the **closest Euclidean distance** to the **N -bit vector** is selected. To reduce the classification error, error correcting codes are used when generating the “code book”

Class	Code Word														
	f_0	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_{10}	f_{11}	f_{12}	f_{13}	f_{14}
0	1	1	0	0	0	0	1	0	1	0	0	1	1	0	1
1	0	0	1	1	1	1	0	1	0	1	1	0	0	1	0
2	1	0	0	1	0	0	0	1	1	1	1	0	1	0	1
3	0	0	1	1	0	1	1	1	0	0	0	0	1	0	1
4	1	1	1	0	1	0	1	1	0	0	1	0	0	0	1
5	0	1	0	0	1	1	0	1	1	1	0	0	0	0	1
6	1	0	1	1	1	0	0	0	0	1	0	1	0	0	1
7	0	0	0	1	1	1	1	0	1	0	1	1	0	0	1
8	1	1	0	1	0	1	1	0	0	1	0	0	0	1	1
9	0	1	1	1	0	0	0	0	1	0	1	0	0	1	1

A code book with $K=9$, $N=15$



STEVENS
INSTITUTE *of* TECHNOLOGY

THE INNOVATION UNIVERSITY®

stevens.edu